

Ridge all concept from scratch to advanced

Ridge Regression: Complete Guide (From Scratch to Advanced)

1. Foundation: Linear Regression & The Problem

1.1 Ordinary Least Squares (OLS)

For linear regression: $y = X\beta + \epsilon$

- OLS minimizes: $RSS(\beta) = ||y - X\beta||^2$
- Solution: $\hat{\beta} = (X^T X)^{-1} X^T y$

1.2 The Problem with OLS

1. **Multicollinearity:** When $X^T X$ is near-singular ($\det \approx 0$)
2. **High variance:** Small changes in data $\rightarrow$ large changes in $\hat{\beta}$
3. **Overfitting:** Especially with many predictors ($p \approx n$ or $p > n$)

2. Ridge Regression Core Concept

2.1 The Intuition

Add **L2 penalty** to control coefficient magnitudes:

text

minimize: $||y - X\beta||^2 + \lambda ||\beta||^2$

where:

- $\lambda \geq 0$ is the **regularization parameter**
- $||\beta||^2 = \sum \beta_j^2$ (excluding β_0 if intercept is separate)

2.2 Geometric Interpretation

- Constrains β to lie within a hypersphere
- Balances fit (RSS) and complexity (coefficient magnitudes)
- **Bias-variance tradeoff:** As λ increases:
 - Bias $\uparrow$, Variance $\downarrow$
 - Coefficients shrink toward zero

3. Mathematical Derivation

3.1 Objective Function

text

$$L(\beta) = (y - X\beta)^T(y - X\beta) + \lambda\beta^T\beta$$

3.2 Closed-Form Solution

text

$$\hat{\beta}_{\text{ridge}} = (X^T X + \lambda I)^{-1} X^T y$$

where I is $p \times p$ identity matrix.

3.3 Solving the Singularity Problem

Even if $X^T X$ is singular, $(X^T X + \lambda I)$ is always invertible for $\lambda > 0$!

Proof:

- $X^T X$ is positive semi-definite: eigenvalues ≥ 0
- $(X^T X + \lambda I)$ has eigenvalues: $(\text{eigen}_i + \lambda) > 0$ for $\lambda > 0$
- All eigenvalues $> 0 \Rightarrow$ matrix is invertible

4. Implementation from Scratch

4.1 Basic Python Implementation

python

```

import numpy as np

class RidgeRegression:
    def __init__(self, alpha=1.0):
        """
        alpha: regularization strength ( $\lambda$ )
        """
        self.alpha = alpha
        self.coef_ = None
        self.intercept_ = None

    def fit(self, X, y):
        # Add intercept
        X_b = np.c_[np.ones((X.shape[0], 1)), X]

        # Identity matrix (don't penalize intercept)
        I = np.eye(X_b.shape[1])
        I[0, 0] = 0 # Don't regularize intercept

        # Ridge solution
        A = X_b.T @ X_b + self.alpha * I
        b = X_b.T @ y
        self.theta = np.linalg.solve(A, b)

        self.intercept_ = self.theta[0]
        self.coef_ = self.theta[1:]

        return self

    def predict(self, X):
        X_b = np.c_[np.ones((X.shape[0], 1)), X]
        return X_b @ self.theta

# Example usage
np.random.seed(42)
n_samples, n_features = 100, 10
X = np.random.randn(n_samples, n_features)
true_coef = np.random.randn(n_features)
y = X @ true_coef + np.random.randn(n_samples) * 0.1

model = RidgeRegression(alpha=1.0)
model.fit(X, y)

```

4.2 With SVD for Numerical Stability

python

```

class RidgeRegressionSVD:
    def __init__(self, alpha=1.0):
        self.alpha = alpha

    def fit(self, X, y):
        # Center data for intercept
        self.X_mean = X.mean(axis=0)
        self.y_mean = y.mean()
        X_centered = X - self.X_mean
        y_centered = y - self.y_mean

        # SVD decomposition
        U, s, Vt = np.linalg.svd(X_centered, full_matrices=False)

        # Ridge solution via SVD
        d = s / (s**2 + self.alpha)
        self.coef_ = Vt.T @ (d * (U.T @ y_centered))
        self.intercept_ = self.y_mean - self.X_mean @ self.coef_

    return self

```

5. Key Properties & Insights

5.1 Bias-Variance Tradeoff

text

$$\text{Var}(\hat{\beta}_{\text{ridge}}) = \sigma^2 V(D^2 / (D^2 + \lambda I)^2) V^T$$

$$\text{Bias}(\hat{\beta}_{\text{ridge}}) = -\lambda V(D^2 + \lambda I)^{-1} V^T \beta$$

where $X = UDV^T$ (SVD decomposition).

5.2 Effective Degrees of Freedom

text

$$\text{df}(\lambda) = \sum (d_j^2 / (d_j^2 + \lambda))$$

where d_j are singular values of X .

- $\lambda = 0 \rightarrow \text{df} = p$ (OLS)
- $\lambda \rightarrow \infty \rightarrow \text{df} \rightarrow 0$ (all coefficients = 0)

5.3 Scale Sensitivity

Ridge is NOT scale invariant!

```
python

# Always standardize features before Ridge!
from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
```

6. Advanced Topics

6.1 Bayesian Interpretation

Ridge $\equiv$ **MAP estimation** with Gaussian priors:

```
text

Prior:  $\beta \sim N(0, \tau^2 I)$ 
Likelihood:  $y|X, \beta \sim N(X\beta, \sigma^2 I)$ 
Posterior:  $\beta|y, X \sim N(\hat{\beta}_{\text{ridge}}, \Sigma)$ 
```

where $\lambda = \sigma^2/\tau^2$

6.2 Kernel Ridge Regression

Extend to non-linear via kernel trick:

```
text

 $\hat{\beta} = (K + \lambda I)^{-1}y$ 
 $f(x) = \sum \alpha_i K(x, x_i)$ 
```

where K is kernel matrix.

```
python

from sklearn.kernel_ridge import KernelRidge
from sklearn.gaussian_process.kernels import RBF

krr = KernelRidge(alpha=1.0, kernel=RBF(length_scale=1.0))
```

6.3 Generalized Ridge

Different penalty per coefficient:

text

$$\min ||y - X\beta||^2 + \sum \lambda_j \beta_j^2$$

Useful when you have prior knowledge about coefficients.

6.4 Ridge for Classification (Ridge Logistic Regression)

text

$$\min \sum [-y_i \log(\sigma(x_i \beta)) - (1-y_i) \log(1-\sigma(x_i \beta))] + \lambda ||\beta||^2$$

where σ is sigmoid function.

7. Practical Considerations

7.1 Choosing λ

Methods:

1. **Cross-validation** (most common):

python

```
from sklearn.linear_model import RidgeCV

model = RidgeCV(alphas=np.logspace(-6, 6, 13),
                 scoring='neg_mean_squared_error',
                 cv=5)
```

2. **Analytical methods:**

- Generalized cross-validation (GCV)
- AIC/BIC (with effective df)

3. **Visualization:** Ridge trace plot

python

```

alphas = np.logspace(-2, 4, 100)
coefs = []
for a in alphas:
    ridge = Ridge(alpha=a)
    ridge.fit(X_scaled, y)
    coefs.append(ridge.coef_)

# Plot coefficient paths
plt.figure(figsize=(10, 6))
for i in range(len(coefs[0])):
    plt.plot(alphas, [c[i] for c in coefs])
plt.xscale('log')
plt.xlabel('λ')
plt.ylabel('Coefficient value')

```

7.2 Comparison with Other Methods

Method	Penalty	Sparsity	When to Use
Ridge	L2	No	Many correlated predictors
Lasso	L1	Yes	Feature selection needed
ElasticNet	L1 + L2	Some	Groups of correlated features

7.3 Feature Selection with Ridge

While Ridge doesn't zero out coefficients:

1. **Strongly shrunk** coefficients $\approx$ less important
2. Use **stability selection**:

```

python

from sklearn.utils import resample

def stability_selection(X, y, n_bootstraps=100, alpha=1.0):
    selected = np.zeros(X.shape[1])
    for _ in range(n_bootstraps):
        X_resample, y_resample = resample(X, y)
        model = Ridge(alpha=alpha)
        model.fit(X_resample, y_resample)
        # Select top k features
        top_k = np.argsort(np.abs(model.coef_))[-5:]

```

```

        selected[top_k] += 1
    return selected / n_bootstraps

```

8. Modern Extensions & Applications

8.1 Ridge for Deep Learning

- **Weight decay** $\equiv$ Ridge regularization
- L2 penalty in neural network loss:

```

python

# PyTorch example
import torch.nn as nn

# Equivalent to Ridge:
optimizer = torch.optim.SGD(model.parameters(),
                              lr=0.01,
                              weight_decay=1e-4) #  $\lambda$ 

```

8.2 Multi-task Ridge

text

$$\min \sum ||Y_k - X\beta_k||^2 + \lambda \sum ||\beta_k||^2$$

Shares information across tasks.

8.3 Time Series: Ridge VAR

Vector AutoRegression with Ridge:

```

python

from sklearn.linear_model import Ridge
import pandas as pd

def ridge_var(data, lags=2, alpha=1.0):
    n_vars = data.shape[1]
    X, Y = [], []
    for t in range(lags, len(data)):
        X.append(data[t-lags:t].flatten())
        Y.append(data[t])

```



```

models = []
for i in range(n_vars):
    model = Ridge(alpha=alpha)
    model.fit(X, [y[i] for y in Y])
    models.append(model)
return models

```

9. Common Pitfalls & Solutions

9.1 Pitfall 1: Forgetting to Standardize

python

```

# WRONG: Without standardization
model.fit(X, y) # Features on different scales

# RIGHT: Always standardize
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
model.fit(X_scaled, y)

```

9.2 Pitfall 2: Intercept Regularization

python

```

# Don't penalize intercept!
class CorrectRidge:
    def fit(self, X, y):
        X_b = np.c_[np.ones((len(X), 1)), X]
        I = np.eye(X_b.shape[1])
        I[0, 0] = 0 # ← Critical!
        # ... rest of implementation

```

9.3 Pitfall 3: λ Range Selection

python

```

# Good practice: Logarithmic grid
alphas = np.logspace(-6, 6, 13)
# Check both very small and very large  $\lambda$ 

```

10. Performance Optimization

10.1 Warm Start for Path Calculation

python

```
def ridge_path(X, y, alphas):
    coefs = np.zeros((len(alphas), X.shape[1]))

    # Start with strongest regularization
    model = Ridge(alpha=alphas[-1], warm_start=True)

    for i, alpha in enumerate(reversed(alphas)):
        model.set_params(alpha=alpha)
        model.fit(X, y)
        coefs[-(i+1)] = model.coef_

    return coefs
```

10.2 Parallel Cross-Validation

python

```
from joblib import Parallel, delayed

def cv_score(alpha, X, y, cv=5):
    scores = []
    for train_idx, val_idx in KFold(cv).split(X):
        X_train, X_val = X[train_idx], X[val_idx]
        y_train, y_val = y[train_idx], y[val_idx]

        model = Ridge(alpha=alpha)
        model.fit(X_train, y_train)
        score = -mean_squared_error(y_val, model.predict(X_val))
        scores.append(score)
    return np.mean(scores)

# Parallel computation
scores = Parallel(n_jobs=-1)(
    delayed(cv_score)(alpha, X, y)
    for alpha in alphas
)
```

11. Case Study: High-Dimensional Genetics Data

python

```
# Example: Gene expression data (p >> n)
n_samples, n_genes = 100, 10000
X = np.random.randn(n_samples, n_genes)
y = np.random.randn(n_samples)

# Standard Ridge struggles with p >> n
# Solution: Pre-selection or Kernel Ridge

# Two-step approach:
# 1. Select top k genes by marginal correlation
correlations = np.abs([np.corrcoef(X[:, i], y)[0, 1]
                       for i in range(n_genes)])
top_genes = np.argsort(correlations)[-100:] # Top 100

# 2. Apply Ridge on selected genes
ridge = RidgeCV(alphas=np.logspace(-4, 4, 9))
ridge.fit(X[:, top_genes], y)
```

12. Future Directions

12.1 Adaptive Ridge

text

$$\min ||y - X\beta||^2 + \lambda \sum w_j \beta_j^2$$

where w_j adaptively chosen (similar to adaptive Lasso).

12.2 Randomized Ridge

For very large p:

python

```
from sklearn.linear_model import Ridge
from sklearn.random_projection import GaussianRandomProjection

# Random projection + Ridge
transformer = GaussianRandomProjection(n_components=100)
X_proj = transformer.fit_transform(X)
ridge = Ridge(alpha=1.0).fit(X_proj, y)
```

Summary

When to use Ridge:

- Multicollinearity present
- All predictors potentially relevant
- $p \approx n$ or $p > n$ but want all features
- Bayesian perspective desired

Key takeaways:

1. Ridge adds L2 penalty → stabilizes coefficients
2. Always standardize features
3. Use cross-validation to choose λ
4. Interpret through Bayesian lens or coefficient paths
5. Modern ML: Ridge is everywhere (weight decay in DL)

Final implementation checklist:

- Standardize features
- Don't regularize intercept
- Use logarithmic grid for λ search
- Validate with cross-validation
- Consider Bayesian interpretation for uncertainty
- Compare with Lasso/ElasticNet for your problem